

DAD

Desenvolvimento de
Aplicações Distribuídas

Message Queues

Agenda

- **Introduction**
- **Programming Model**
- **Case-study:**
 - Websphere MQ
 - Java Messaging Service

Introduction

■ **Indirect Communication**

- ✚ loose-coupling
- ✚ queues are first-class entities
- ✚ no explicit end-point address of senders/receivers

■ ***point-to-point service***

- ✚ not one-to-many as groups and pub-sub

■ **used heavily in:**

- ✚ Enterprise Application Integration (EAI) scenarios
- ✚ *commercial transaction processing systems*

Programming Model

■ What is a Queue:

- ✚ communication end-point
- ✚ contents:
 - ✚ ordered set of persistent messages
 - ✚ multiple producers and consumers
 - ✚ messages only consumed by one process

Programming Model

■ **Operations on Queues:**

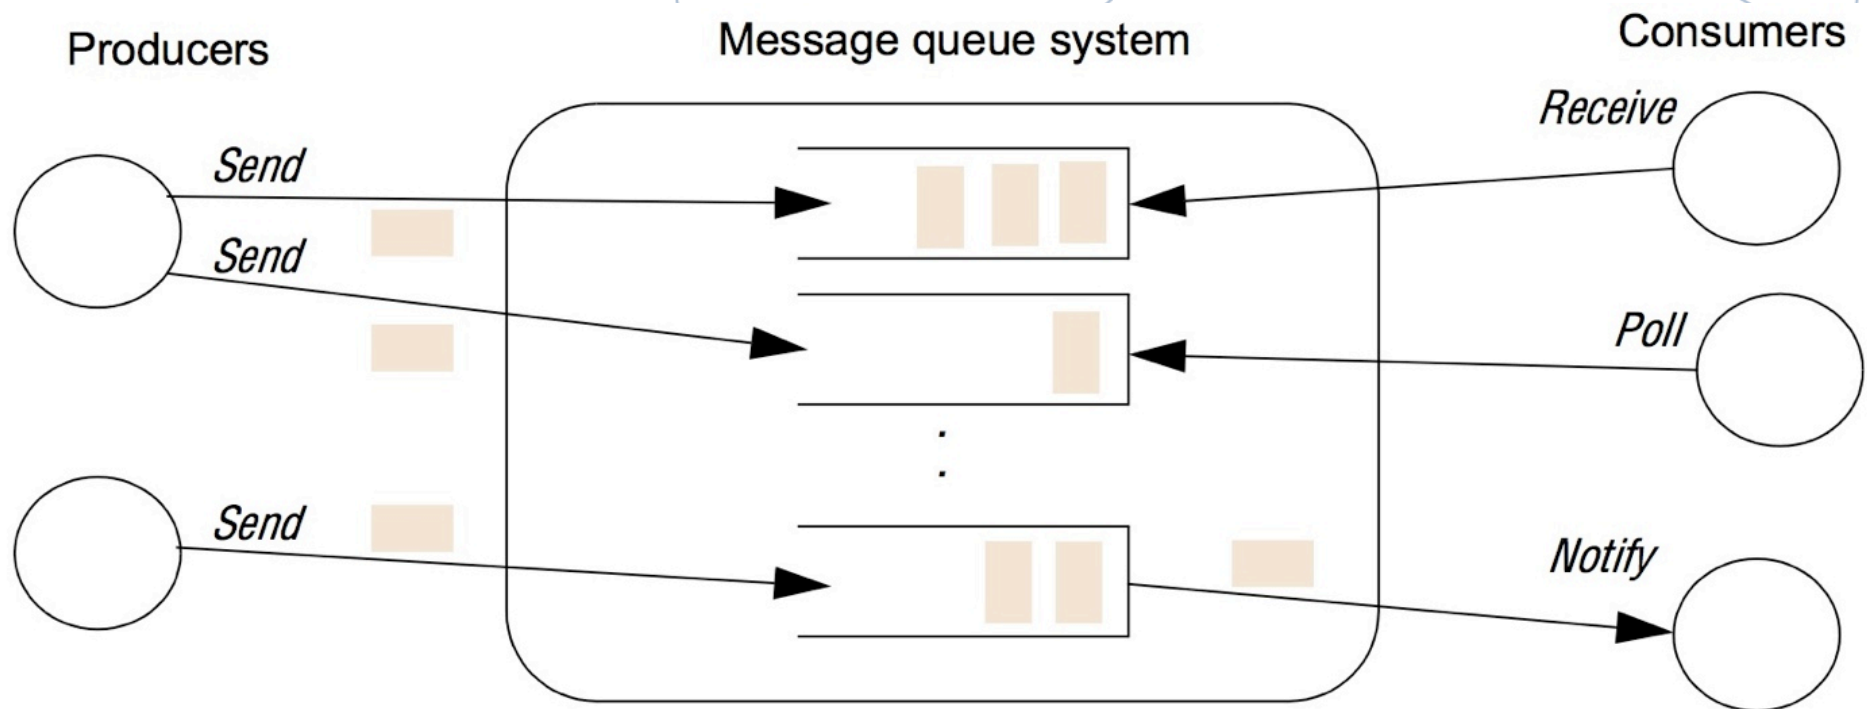
- ✚ send (q, msg, ...)
- ✚ msg = receive (q,...)
- ✚ possibly integrated into transactions
 - ✳ atomicity of all operations and messages sent

■ **Receive semantics:**

- ✚ blocking (receive)
 - ✳ until appropriate message available
- ✚ non-blocking (poll)
 - ✳ return message if available, otherwise return *not_available*
- ✚ notify
 - ✳ raise event when message is available in queue

Programming Model

- The message queue paradigm**



Programming Model

■ **Queueing policy:**

- ✚ order by which messages are consumed
 - ✧ FIFO, most common
 - ✧ priority-based
 - ✧ based on message properties

■ **Message structure:**

- ✚ destination (queue)
- ✚ metadata (e.g., priority, delivery mode)
- ✚ body (opaque to system)
 - ✧ can be a row or set of rows in a database

Properties

■ **Reliable Delivery:**

- ✚ messages are persistent
 - ✦ stored on disk indefinitely until consumed
 - ✦ including transactional support, journalling, logging, ...

■ **Validity:**

- ✚ any message sent is eventually received

■ **Integrity:**

- ✚ message received identical to the one sent
- ✚ no message is delivered twice

Websphere MQ (IBM middleware)

■ ***queue managers***

- ✚ host and manage queues, enforce reliability

■ **queue operations**


- ✚ defined by a Message Queue Interface (MQI)
 - ✦ connect /disconnect (MQCONN / MQDISC)
 - ✦ send message (MQPUT)
 - ✦ receive message (MQGET)

■ **channels:**

- ✚ asynchronous message channels between *queue managers*
 - ✦ unidirectional, managed by 2 message channel agents (MCA)
- ✚ client channels (senders/receivers to queue managers)

Websphere MQ

- **simple networked topology in Websphere MQ**



Websphere MQ

■ ***Networks of Queue Managers:***

- ✚ queue managers linked in federated structure
 - ✧ with message channels
 - ✧ routing tables at each queue manager
- ✚ arbitrary topologies possible: trees, meshes, bus-based...

■ ***Hub-and-Spoke Approach:***

- ✚ one QM is the hub
 - ✧ hosts most services, executed upon message reception
- ✚ other, several QM are designated as spokes (*relays*)
 - ✧ placed around the network for geo-coverage and load-balancing
 - ✧ clients connect only to (near-by) spokes
- ✚ widely used but hub potential bottleneck and single-failure
 - ✧ use *queue manager clusters* (for QM replication)

JMS: Java Messaging Service

■ **Java standard specification**

- ✚ indirect communication
- ✚ attempt to partial unify pub-sub and message queues
- ✚ many implementations;
 - ✦ Joram/OW2, JBoss, Sun OpenMQ, Apache ActiveMQ, OpenJMS
 - ✦ also Wepshere MQ.

■ **Key roles:**

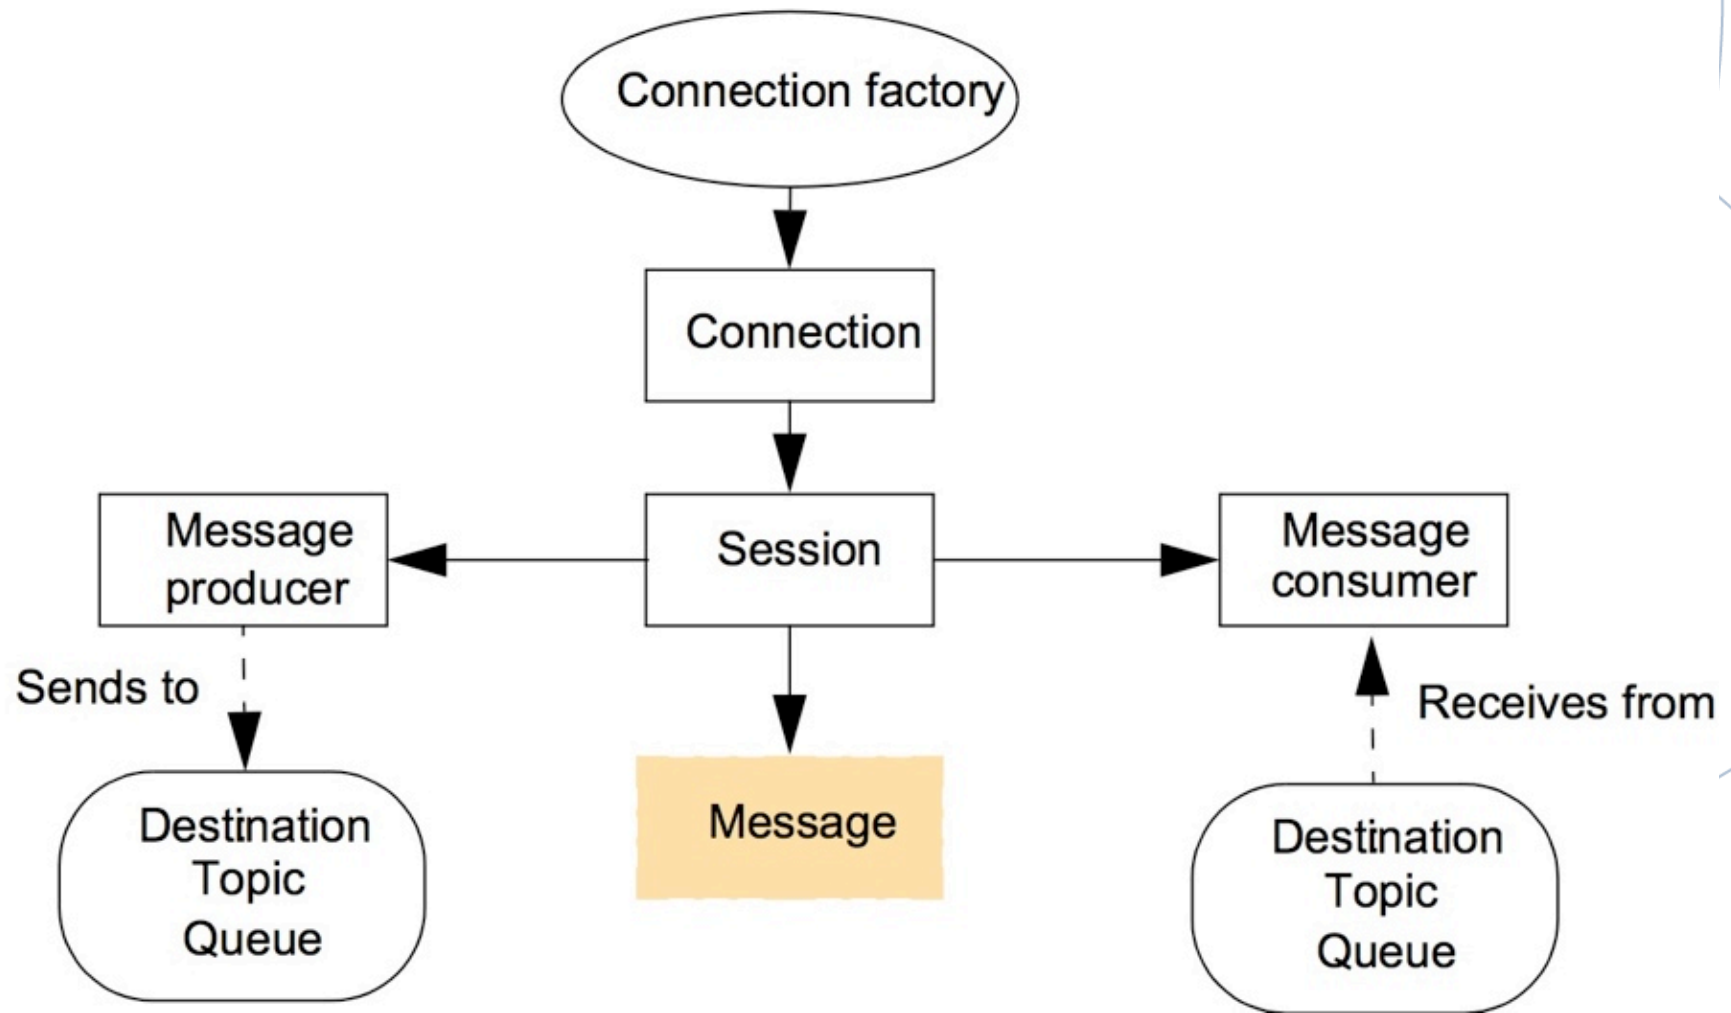
- ✚ client (sender/receiver program)
- ✚ provider (system that implements JMS, e.g., a server)
- ✚ message (an object sent to the queue).

■ **Programming model:**

- ✚ Connections, Sessions, Topics, Queues, Transactions

JMS: Java Messaging Service

- The programming model offered by JMS



JMS: class *FireAlarmJMS*

```

import javax.jms.*;
import javax.naming.*;
public class FireAlarmJMS {

    public void raise() {
        try {
            Context ctx = new InitialContext();
            TopicConnectionFactory topicFactory =
            (TopicConnectionFactory)ctx.lookup ("TopicConnectionFactory");
            Topic topic = (Topic)ctx.lookup("Alarms");
            TopicConnection topicConn =
                topicConnectionFactory.createTopicConnection();
            TopicSession topicSess = topicConn.createTopicSession(false,
                Session.AUTO_ACKNOWLEDGE);
            TopicPublisher topicPub = topicSess.createPublisher(topic);
            TextMessage msg = topicSess.createTextMessage();
            msg.setText("Fire!");
            topicPub.publish(message);
        } catch (Exception e) {
        }
    }
  
```

JMS: class *FireAlarmConsumerJMS*

```

import javax.jms.*; import javax.naming.*;
public class FireAlarmConsumerJMS
public String await() {
    try {
        Context ctx = new InitialContext();
        TopicConnectionFactory topicFactory =
            (TopicConnectionFactory)ctx.lookup("TopicConnectionFactory");
        Topic topic = (Topic)ctx.lookup("Alarms");
        TopicConnection topicConn =
            topicConnectionFactory.createTopicConnection();
        TopicSession topicSess = topicConn.createTopicSession(false,
            Session.AUTO_ACKNOWLEDGE);
        TopicSubscriber topicSub = topicSess.createSubscriber(topic);
        topicSub.start();
        TextMessage msg = (TextMessage) topicSub.receive();
        return msg.getText();
    } catch (Exception e) {
        return null;
    }
}

```

Summary

- **Introduction**
- **Programming Model**
- **Case-study:**
 - Websphere MQ
 - Java Messaging Service

DAD
Desenvolvimento de
Aplicações Distribuídas

Shared Memory Approaches

Agenda

- **Introduction**
- **Distributed Shared Memory**
- **Tuple-space Communication**
 - ✚ other approaches: replication, partitioning
- **Case-study:**
 - ✚ JavaSpaces
- **Indirect Communication Summary**

Introduction

■ previous Indirect Communication based on *message-passing*

- ✚ *group communication*
- ✚ *publish-subscribe*
- ✚ *message queues*

■ Shared Memory approaches

- ✚ abstraction of shared address/data space
- ✚ accessed with reading and writing operations
- ✚ located by *memory address* or by *content (associative)*

Distributed Shared Memory (DSM)

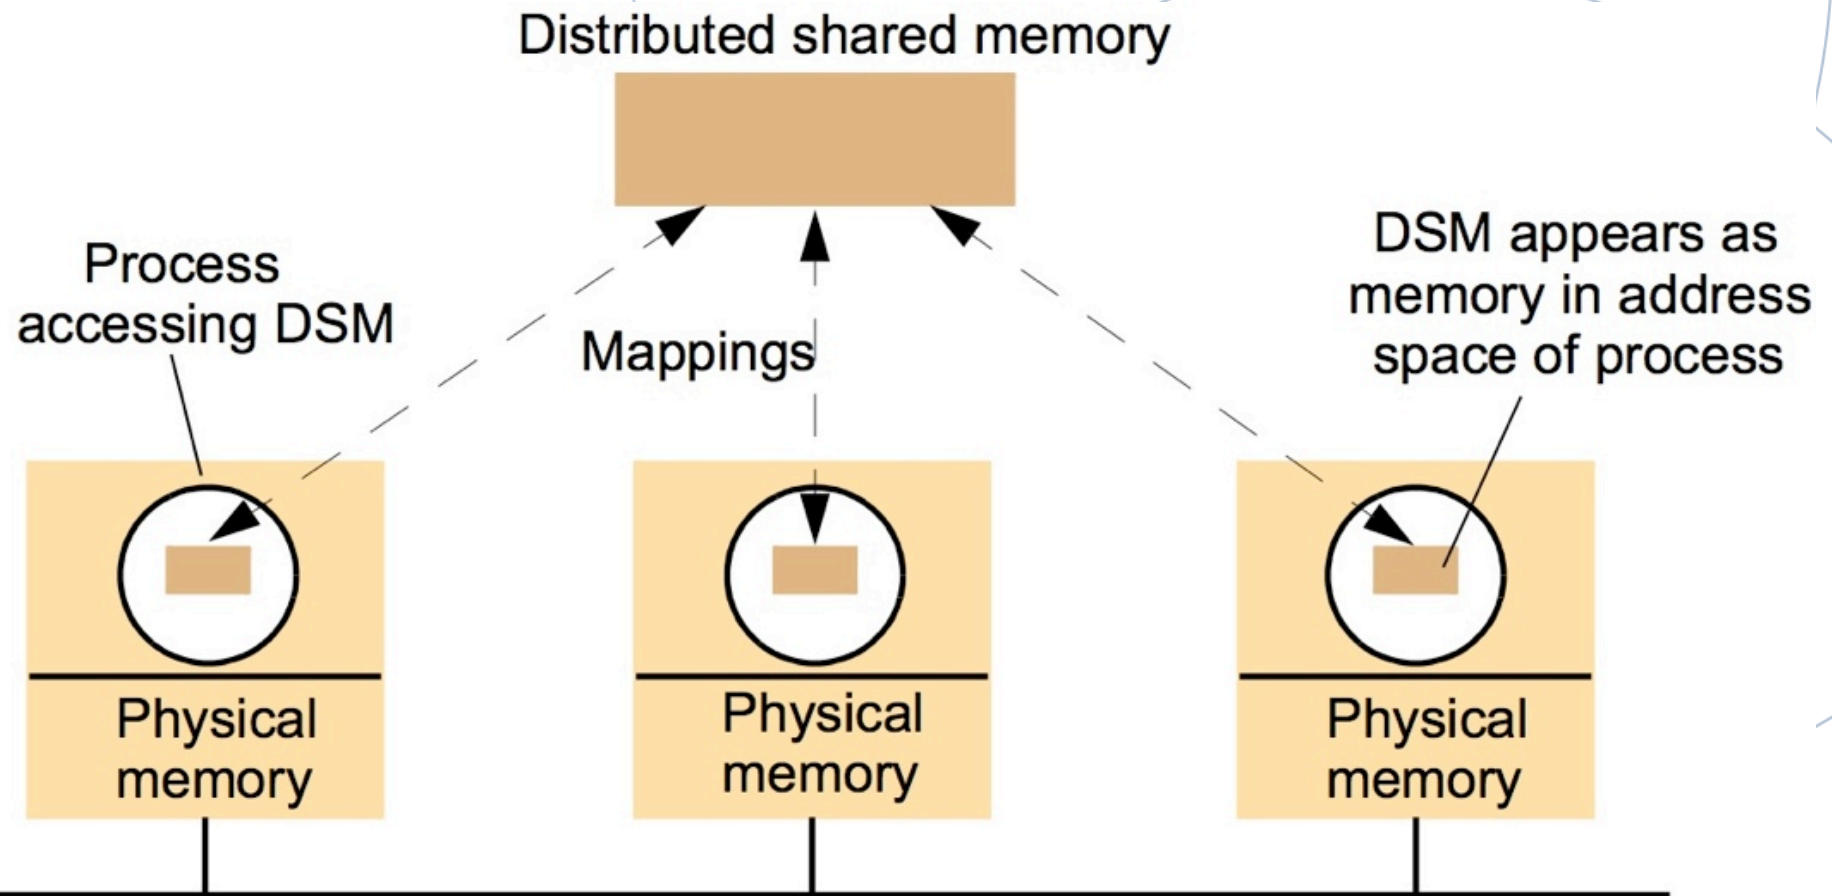
■ **abstraction for sharing data**

- ✚ between computers with separate physical memory
- ✚ shared memory regions
 - ✧ mapped in each process as *ordinary* memory in address space
- ✚ data accessed/modified by read/write CPU instructions
- ✚ supported transparently by runtime system

■ **DSM processes manipulate shared logical global address space**

- ✚ while in fact physical memory is distributed

Distributed Shared Memory



Distributed Shared Memory (DSM)

■ Main Goal

- ✚ free programmers from concerns with message-passing
- ✚ common use in parallel and distributed applications
 - ✳ accessing shared data items individually and directly (i.e., variables, objects, arrays, ...)
 - ✳ easier to express parallel algorithms
 - ✳ not appropriate for client-server, lacks modularity and protection

■ Implementation Issues

- ✚ message passing cannot be avoided in distribution
- ✚ DSM runtime has to exchange updates among processes
- ✚ DSM systems manage local copies of replicated data for access speed
- ✚ evolution: shared-memory multiprocessors, non-uniform memory access (NUMA), distributed-memory multiprocessors, clusters

Tuple Space Communication

■ **abstraction for sharing data**

- ✚ unlike DSM, no direct addressing,
- ✚ data located by pattern matching on content
 - ✱ early form of content-addressable memory
- ✚ organized as *tuples* (akin to database records)
- ✚ tuples retrieved and created, never modified
 - ✱ makes synchronization easier/unnecessary

■ **relevant examples**

- ✚ Linda, IBM Tspaces, Sun JavaSpaces

Tuple Space Communication

■ Programming Model

✚ data space

- ✳ shared collection of tuples (sequence of typed data fields):
 - e.g., <“fred”, 1958>, <“sid”, 1964>, <4, 9.8, “yes”>
- ✳ processes shared data by accessing the same tuple space

✚ operations

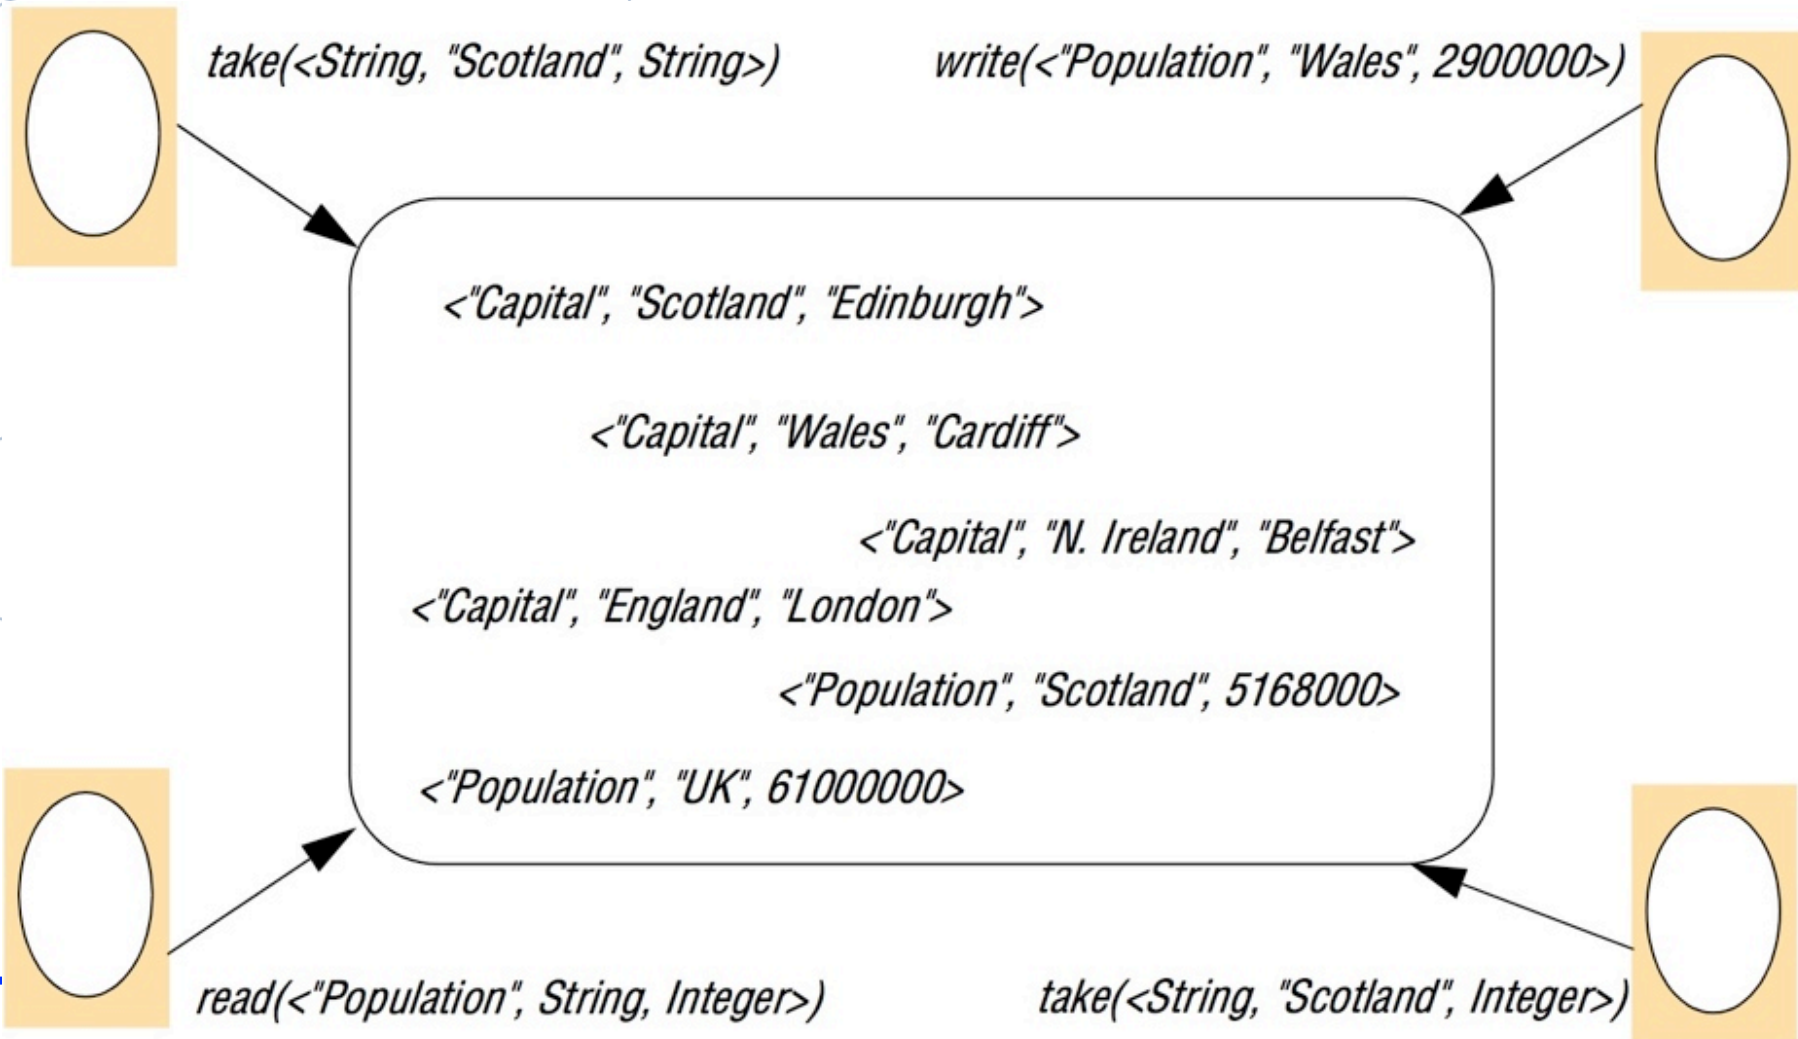
- ✳ **write (out)**: creates and adds tuples to the tuple space
 - without affecting any other tuple
- ✳ **read (rd)**: retrieves content of one tuple
 - without affecting any other tuple, tuple may be read several times
- ✳ **take (in)**: retrieves one tuple and removes it from the tuple space
 - each tuple is *consumed* by only one *take* operation

✚ tuple specification (e.g., <*String*, “Scotland”, *String*>

- ✳ *template* to find *associatively*, any matching tuple in the space

Tuple Space Communication

Programming model and examples of tuple operations



Tuple Space Communication

■ Synchronization

- ✚ read and take operations are blocking
 - ✱ until there is matching tuple in the space
- ✚ no direct access to tuples in the tuple space
 - ✱ processes have to replace tuples instead of modifying them
- ✚ e.g.,
 - ✱ shared counter <“counter”, 64> manipulated by several processes
 - ✱ to increment counter, a process performs
 - ✱ `<s, count> := myTS.take(<“counter”, integer>);`
 - ✱ `myTS.write(<“counter”, count+1>);`

Tuple Space Communication

■ Important Properties:

✚ space uncoupling

- ✧ tuple in space may originate from any process and may be delivered to any number of processes

✚ time uncoupling

- ✧ tuple placed in the tuple space will remain until removed (potentially indefinitely), hence sender and receiver need not overlap in time

■ Variations

- ✚ multiple tuple spaces, per-user, per-application, system
- ✚ distributed implementations (original ones were centralized)
- ✚ objects (methods, attributes) stored as tuple fields

Tuple Space (Other Approaches)

■ Replicated Tuple Spaces

- ✚ goal: fault-tolerance and scalability

■ State-machine approach

- ✚ assume tuple space behaves as state-machine
- ✚ to ensure consistency, replicas must:
 - ✚ 1. begin with the same data (empty tuple space)
 - ✚ 2. execute events in the same order
 - ✚ 3. react deterministically to each event
- ✚ property (2) can be ensured by total order multicast (coming soon!)

Tuple Space (Other Approaches)

■ Replicated Tuple Spaces

- ✚ goal: fault-tolerance and scalability

■ Xu and Liskov

- ✚ optimizes replication strategy, and uses partitioning
- ✚ leverages semantics of particular tuple space operations
- ✚ processes:
 - ✱ set of workers that perform computation on the tuple space
 - ✱ set of tuple space replicas
 - ✱ worker and replica may be combined in a single physical node
- ✚ network assumed to fail:
 - ✱ may lose, duplicate, delay messages or delivered them out of order
- ✚ operations executed in the context of current view
- ✚ **view**: agreed set of replicas of the tuple space

Replicated Tuple Spaces: Xu and Liskov (operation algorithms)

■ **Write:**

- ✚ 1. The requesting site multicasts the write request to all members of the view;
- ✚ 2. On receiving this request, members insert the tuple into their replica and acknowledge this action;
- ✚ 3. Step 1 is repeated until all acknowledgements are received

Replicated Tuple Spaces: Xu and Liskov (operation algorithms)

■ Read

- ✚ 1. The requesting site multicasts the read request to all members of the view;
- ✚ 2. On receiving this request, a member returns a matching tuple to the requestor;
- ✚ 3. The requestor returns the first matching tuple received as the result of the operation (ignoring others);
- ✚ 4. Step 1 is repeated until at least one response is received.

Replicated Tuple Spaces: Xu and Liskov (operation algorithms)

■ **Take (*Phase 1: Selecting the tuple to be removed*)**

- ✚ 1. Requesting site multicasts the take request to all members of view;
- ✚ 2. On receiving this request, each replica acquires a lock on the associated tuple set and, if the lock cannot be acquired, the take request is rejected;
- ✚ 3. All accepting members reply with the set of all matching tuples;
- ✚ 4. Step 1 is repeated until all sites have accepted the request and responded with their set of tuples and the intersection is non-null;
- ✚ 5. A particular tuple is selected as the result of the operation (selected randomly from the intersection of all the replies);
- ✚ 6. If only a minority accept the request, this minority are asked to release their locks and phase 1 repeats.

Replicated Tuple Spaces: Xu and Liskov (operation algorithms)

■ **Take (*Phase2: Removing the selected tuple*)**

- ✚ 1. The requesting site multicasts a remove request to all members of the view citing the tuple to be removed;
- ✚ 2. On receiving this request, members remove the tuple from their replica, send an acknowledgement and release the lock;
- ✚ 3. Step 1 is repeated until all acknowledgements are received.

Replicated Tuple Spaces: Xu and Liskov (operation latency)

- Minimize delay given the **semantics** of the three tuple space operations:
 - ✚ read operations only block until the first replica responds to the request
 - ✚ write operations can return immediately.
 - ✚ take operations block until the end of phase 1, when the tuple to be deleted has been agreed.

Replicated Tuple Spaces: Xu and Liskov

■ Issues with concurrency:

- ✚ e.g., **read** operation may access tuple that should have been deleted in the 2nd phase of **take** operation

■ Solution: enforce additional constraints

- ✚ operations of each worker must be executed at each replica in the same order they were issued by worker
 - ✚ FIFO order (or client order)

Replicated Tuple Spaces: Xu and Liskov

■ **Issues with concurrency:**

- ✚ e.g., **read** operation may access tuple that should have been deleted in the 2nd phase of **take** operation

■ **Solution: enforce additional constraints**

- ✚ operations of each worker must be executed at each replica in the same order they were issued by worker
 - ✚ FIFO order (or client order)
- ✚ a **write** operation must not be executed by any replica until
 - ✚ all previous **take** operations issued by the same worker have completed at all replicas in the worker's *view*.

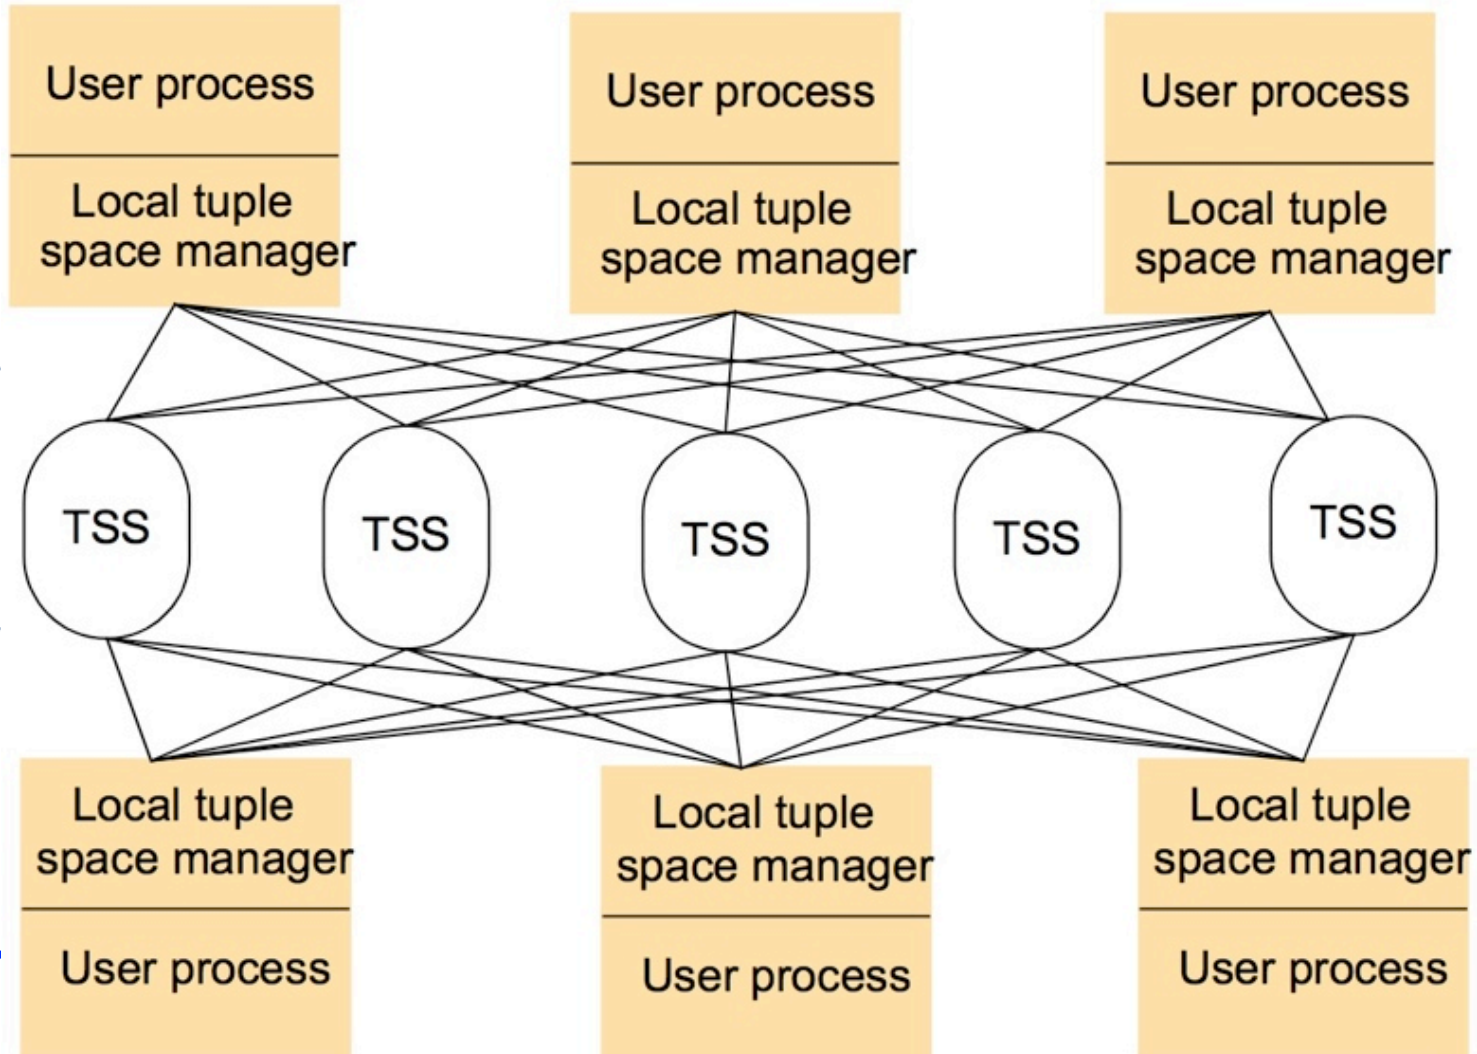
Partitioned Tuple Spaces

■ **Partitioned Tuple Spaces**

- ✚ Linda kernel at Univ. York
- ✚ tuples partitioned across Tuple Space Servers (TSS)
- ✚ no replication, sole goal to increase performance
- ✚ tuple specifications specify type or values
- ✚ tuple location decided by hashing algorithm
 - ✳ generates set of possible servers that may contain matching tuples
 - ✳ local linear searching
- ✚ e.g., of peer-to-peer implementation: PeerSpaces

Partitioned Tuple Spaces

■ Architecture of Linda Kernel Univ. York



JavaSpaces (Jini / Apache River)

■ Programming Model

<i>Operation</i>	<i>Effect</i>
<i>Lease write(Entry e, Transaction txn, long lease)</i>	Places an entry into a particular JavaSpace
<i>Entry read(Entry tmpl, Transaction txn, long timeout)</i>	Returns a copy of an entry matching a specified template
<i>Entry readIfExists(Entry tmpl, Transaction txn, long timeout)</i>	As above, but not blocking
<i>Entry take(Entry tmpl, Transaction txn, long timeout)</i>	Retrieves (and removes) an entry matching a specified template
<i>Entry takeIfExists(Entry tmpl, Transaction txn, long timeout)</i>	As above, but not blocking
<i>EventRegistration notify(Entry tmpl, Transaction txn, RemoteEventListener listen, long lease, MarshalledObject handback)</i>	Notifies a process if a tuple matching a specified template is written to a JavaSpace

JavaSpaces example

```
import net.jini.core.entry.*;  
public class AlarmTupleJS implements Entry {  
    public String alarmType;  
    public AlarmTupleJS() { }  
}  
    public AlarmTupleJS(String alarmType) {  
        this.alarmType = alarmType;  
    }  
}
```


JavaSpaces example

```
import net.jini.space.JavaSpace;  
public class FireAlarmJS {  
    public void raise() {  
        try {  
            JavaSpace space = SpaceAccessor.findSpace("AlarmSpace");  
            AlarmTupleJS tuple = new AlarmTupleJS("Fire!");  
            space.write(tuple, null, 60*60*1000);  
        catch (Exception e) {  
        }  
    }  
}
```

JavaSpaces example

```

import net.jini.space.JavaSpace;
public class FireAlarmConsumerJS {
    public String await() {
        try {
            JavaSpace space = SpaceAccessor.findSpace();
            AlarmTupleJS template = new AlarmTupleJS("Fire!");
            AlarmTupleJS recvd = (AlarmTupleJS) space.read(template, null,
                Long.MAX_VALUE);
            return recvd.alarmType;
        }
        catch (Exception e) {
            return null;
        }
    }
}

```

Indirect Communication (Summary)

	<i>Groups</i>	<i>Publish-subscribe systems</i>	<i>Message queues</i>	<i>DSM</i>	<i>Tuple spaces</i>
<i>Space-uncoupled</i>	Yes	Yes	Yes	Yes	Yes
<i>Time-uncoupled</i>	Possible	Possible	Yes	Yes	Yes
<i>Style of service</i>	Communication-based	Communication-based	Communication-based	State-based	State-based
<i>Communication pattern</i>	1-to-many	1-to-many	1-to-1	1-to-many	1-1 or 1-to-many
<i>Main intent</i>	Reliable distributed computing	Information dissemination or EAI; mobile and ubiquitous systems	Information dissemination or EAI; commercial transaction processing	Parallel and distributed computation	Parallel and distributed computation; mobile and ubiquitous systems
<i>Scalability</i>	Limited	Possible	Possible	Limited	Limited
<i>Associative</i>	No	Content-based publish-subscribe only	No	No	Yes

Summary

- **Introduction**
- **Distributed Shared Memory**
- **Tuple-space Communication**
- **Case-study:**
 - JavaSpaces
- **Indirect Communication Summary**